Sweatshop on P-Chess: lane `core`, rework of PR #10

Corrupt stats files and a repertoire check skipped during the catalog load (HIST-001, OPEN-002, OPEN-004); GPT-6.1 Sol lineup

Foreman, session `sweatshop/core/2026-10-25-0159`

25 October 2026

Summary

The independent review of session PR #10 at head `6fbac3e` (`gpt-6.1-sol xhigh`, run by the supervising session) returned "do not merge" with two blocking findings. The foreman reopened the three owner tickets on the session (`24783fc`) and relaunched the same session at 11:01, `-Models 'stage 2 gpt-6.1-sol high, hard gpt-6.1-sol xhigh, stage 3 gpt-6.1-sol xhigh'`, pid 17404. The driver ended at 11:43 with "Nothing runnable". This report continues `2026-10-25-sweatshop-day-core.pdf`, which covers the run that built the ten tickets.

- **Produced:** the three reopened tickets repaired test first, re-reviewed and merged again on the session: HIST-001, OPEN-002, OPEN-004, each with no new finding. PR #10 still holds 10 tickets, all `Review: agent`, `Verdict: Approve`. Gate 241 passed, 0 failed (from 235).
- **Cost: $4.10** at API list price over 6 stages: `gpt-6.1-sol xhigh` $3.63 ($1.55 as hard implementer, $2.08 in review), `gpt-6.1-sol high` $0.48. Codex quota: 3 points of the weekly window (`gpt-6.1-sol xhigh` 2, outside the stages 1, `gpt-6.1-sol high` 0).
- **What went wrong:** the three defects were present when each ticket was first approved; the ticket reviews did not try a hand-edited or partly written stats file, or a setting changed while the catalog loads. The session review caught them.
- **Provisional verdict:** a cheap, clean rework ($1.37 per ticket, 11–15 min per ticket end to end). The lesson is for the ticket reviews: persisted files and settings changed during an `await` need a probe.

What was done

HIST-001: a stats file that parses but holds a partial result record (only `wins`) or a bare number no longer crashes the end of the game. The bad record counts as zeros, every valid record is kept, and the result and review screens show a warning. `core/history.gd`. Before, the game-end accounting raised a script error reading `losses`, and the review opened as if nothing had happened.

OPEN-002: a practice count saved as text, or negative, is read as 0 and the other counts are kept; Maia's book still picks its practice target, and the result screen shows the same warning. `core/history.gd` (`practice_counts`). Before, the book's target selection raised a script error.

OPEN-004: turning on "Check my moves against my repertoire" while the opening catalog is loading no longer starts a game with an empty repertoire. After the catalog load, `ui/main.gd` decides again whether the repertoire is needed and, if its load failed, returns to the setup screen with the error. Before, the decision was taken before the wait, and the review then showed no repertoire badges.

Timeline

10:55	Supervisor relays the PR #10 review (do not merge, 2 blocking). Foreman: owners from `git blame` (HIST-001 `7e5b8a4`, OPEN-002 `cf0419f`, OPEN-004 `cf5a29a`); tickets reopened with the findings, regressions and HIST-001's second-reopen diagnosis (`24783fc`).
11:01	Relaunch, pid 17404, same session.
11:16	HIST-001 (hard) merged again, 0 findings.
11:30	OPEN-002 (hard) merged again, 0 findings.
11:42	OPEN-004 merged again, 0 findings; "Nothing runnable".

Per stage

When	Ticket	Stage	Model	Min	In	Cache	Out	Cost	Outcome
11:11	HIST-001	impl	sol xhigh	10	2.3M	97%	15k	0.73	to-review
11:16	HIST-001	review	sol xhigh	5	1.5M	91%	10k	0.62	merged
11:25	OPEN-002	impl	sol xhigh	8	2.6M	97%	15k	0.82	to-review
11:30	OPEN-002	review	sol xhigh	6	1.9M	91%	11k	0.80	merged
11:38	OPEN-004	impl	sol high	7	1.4M	95%	8k	0.48	to-review
11:42	OPEN-004	review	sol xhigh	5	1.6M	92%	10k	0.65	merged
				41	11.3M		69k	4.10	

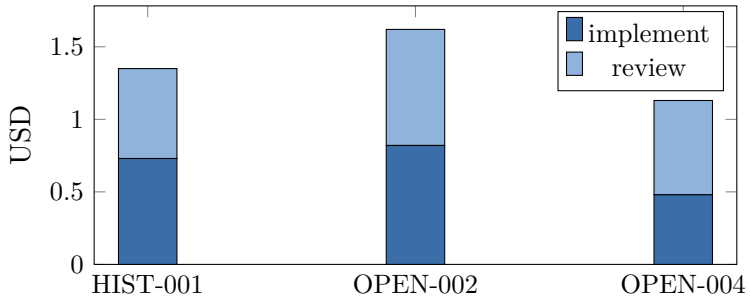

Comparison with previous runs

From `run-log.md`: the run that built these tickets cost $25.44 for 10; this rework adds $4.10, so the PR's ten tickets cost $29.54 ($2.95 each). The previous session-review rework of this lane (UI-004 and UI-005 on Opus, PR #8) cost $8.28 for two tickets.

Observations per model

- **gpt-6.1-sol xhigh, hard stage 2 (HIST-001, OPEN-002)**: red tests first through the real Main (result screen, review), then the smallest validation in `core/history.gd`; both show the corruption as a warning on the result screen.
- **gpt-6.1-sol high, stage 2 (OPEN-004)**: one regression through the real Main, Setup and Settings, then the re-check after the `await`; 7 min.
- **gpt-6.1-sol xhigh, review**: three approvals, no code in review (no round-0 rows).

Driver and runtime issues

None. The reopen kept earlier `Verdict: Reopen` lines unquoted so the driver counted HIST-001's two reopens and gave it the hard implementer; every other line of the closed records was quoted so the driver read no stale receipt or verdict.

Recommendations

1. Ticket reviews of persistence code should load a hand-edited file (a missing key, a wrong type) before approving. All three findings of this session review would have been caught that way or by changing a setting during an `await`.
2. Keep the gpt-6.1-sol lineup: scoreboard row `gpt-6.1-sol high` → `gpt-6.1-sol xhigh` 105 reviews, 17% reopened, S2 rate 16%.

Human debt and proxy decisions

No `Débito humano:` lines and no new `Proxy decided` lines in this rework; the three tickets' earlier proxy decisions (listed in the previous report) stand unchanged.

Reopen attribution

Session-review discoveries, recorded with `Discovery: session-review` and `Round: -`; they do not enter the reopen rate. Judged by the foreman (Opus); the reviewers are `gpt-6.1-sol`.

Ticket	Finding	Label	Why
HIST-001	Partial or scalar result record in `stats.cfg`: script error, no warning	S2-explícito/ código	Criteria 3 and 5. Review miss: yes.
OPEN-002	Non-integer practice count: script error in target selection	S2-implícito	No criterion names corrupt input; failure path of its own persistence. Review miss: yes.
OPEN-004	Checking enabled during the catalog load skips the required load check	S2-explícito/ código	Criterion 2. Review miss: yes.

Scoreboard

Same method as the previous report (lane clone for P-Chess, other repos' main checkouts). Rows that moved: `gpt-6.1-sol high` → `gpt-6.1-sol xhigh` 102 → 105 reviews, 16 → 18 reopened (17%), S2 rate 16%; `gpt-6.1-sol xhigh` → `gpt-6.1-sol xhigh` 4 → 10 reviews, 2 reopened (20%, an anecdote); session-review discoveries 13 → 26 across repos (3 here, the rest EasyMath's), 20 of 22 known were earlier review misses; routine approvals 14 → 17.

Pending for the human

1. Re-review and merge PR #10 (the supervising session runs the re-review).
2. The manual checks and stage-1 triage listed in the previous report are unchanged.